




Tiny Recursive Model (TRM) — attention-free MLP-mixer Sudoku variant. A single shared 2-layer net f is applied $T \times (n+1)$ times per improvement step (stock $T=3, n=6$; our fast config $T=2, n=3$), deep-supervised for up to $N_{\text{sup}}=16$ steps. $\sim 5M$ params.