


FFN-slot phase breakdown (batch 48 x 512 = 24,576 tokens): paper exp_n_0127 vs the int8 read, at both matmul precisions



NOT GEOMETRY-MATCHED: the int8 rows read TWO scored cells per table from a 256-cell table (nap 8), exp_n_0127 reads ONE hard cell per table from a 128-cell table (nap 7), so the int8 rows do strictly more work per token.

All "today" rows, including BEFORE / AFTER, measured in one interleave per precision (3 fresh launches, mean of per-launch medians); every stage segment is a measured stage timing, none inferred.